

BUY, KEEP, OR SELL: ECONOMIC GROWTH AND THE MARKET FOR IDEAS

BY UFUK AKCIGIT, MURAT ALP CELIK, AND JEREMY GREENWOOD¹

An endogenous growth model is developed where each period firms invest in researching and developing new ideas. An idea increases a firm's productivity. By how much depends on the technological propinquity between an idea and the firm's line of business. Ideas can be bought and sold on a market for patents. A firm can sell an idea that is not relevant to its business or buy one if it fails to innovate. The developed model is matched up with stylized facts about the market for patents in the United States. The analysis gauges how efficiency in the patent market affects growth.

KEYWORDS: Growth, ideas, innovation, misallocation, patents, patent agents, research and development, search frictions, technological propinquity.

1. INTRODUCTION

NEW IDEAS ARE THE SEEDS for economic growth. Rising living standards depend on the effectiveness of transforming new ideas into consumer products or production processes. Incarnating an idea into a product or a production process is by no means immediate. Someone must have a vision or an application for the idea and the know-how to implement it. These are often people who work in areas related to the end-use of an idea.

For example, in 1849 Walter Hunt was granted a patent for the safety pin. In the abstract for the patent, Walter Hunt wrote “(t)he distinguishing feature of the invention consist in the construction of a pin made of a piece of wire or metal combining a spring, and a clasp or catch, in which catch the point of the said pin is forced and by its own spring securely retained”—see Figure 1 for his patent application.² Hunt was a mechanic by trade and filed patents for various things, such as ice boats, machines for cutting nails, and repeating guns. What is interesting about this innovation is that Hunt sold his patent to W. R. Grace and Company for about \$10,000 (in today's dollars). W. R. Grace and Company mass-produced the safety pin and made millions.

¹The authors are indebted to Douglas Hanley for help with the patent data. They also thank Jess Benhabib, Sam Kortum, Yunan Li, Laura Liu, Carlos Serrano, Stefanie Stantcheva, and Randy Wright for comments. The authors are also grateful to four referees for their suggestions. Pengfei Han and David Zarruk provided excellent research assistance. Wes Cohen and Mike Roach kindly supplied the authors with data on licensing from the Carnegie Mellon Survey. Akcigit gratefully acknowledges financial support from the Ewing M. Kauffman Foundation and from the Alfred P. Sloan Foundation.

²Patents are publicly disclosed and filed at the United States Patent and Trademark Office. Each patent application has a full description of the invention and drawings to illustrate the embodiments.

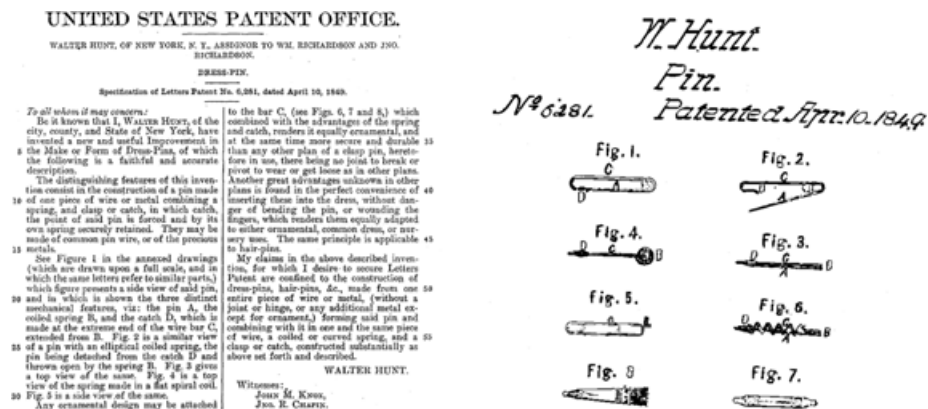


FIGURE 1.—Patent US6281. Walter Hunt's patent for the safety pin.

Walter Hunt by no means was an exception. Firms often develop patents that are not close to their primary business activity.³ Recently released data on the U.S. market for patents indicate that a large fraction of patents are sold by firms, which developed the ideas, to other firms. Specifically, among all the patents registered between 1976 and 2006 in the United States Patent and Trademark Office (USPTO), 16% are traded and this number goes up to 20% among domestic patents.⁴ For economic progress, not only the possibility of exchange, but also the speed of that process is important. USPTO data show that new patents are sold among firms on average within 5.48 years (with a standard deviation of 4.58 years).

An analysis of the patent data in Section 3 uncovers some important facts about the nature of these exchanges. A notion of technological propinquity between a patent and a firm is developed. The key findings are:

1. A patent contributes more to a firm's stock market value if it is closer to the firm in terms of technological distance.
2. A patent is more likely to be sold the more distant it is to the inventing firm.
3. A patent is technologically closer to the buying firm than to the selling firm.

The above observations raise important questions that have been left unanswered by the existing literature: How sizeable is the misallocation of ideas across firms? How does efficiency in the market for ideas affect economic

³Some background material on this was presented in Akcigit, Celik, and Greenwood (2015, Appendix 10).

⁴These numbers do not include patent transfers due to mergers and acquisitions (M&A) and licensing. They include only firm-to-firm patent transfers and exclude within-firm patent transfers as well as patents sold by individuals. See Empirical Appendix B of the Supplemental Material (Akcigit, Celik, and Greenwood (2016)) for data construction.

growth? Do frictions in the market for ideas lead to more in-house R&D or do they discourage innovation overall? This paper attempts to answer these questions.

1.1. *The Analysis*

To analyze the impact that a market for patents has on the macroeconomy, a search-theoretic growth model is built. The framework is developed in Section 2. Each period, firms invest in research and development (R&D). Sometimes this process generates an idea; other times it does not. Each firm operates within a particular technology class, which is fixed over time. An idea increases a firm's productivity. In the current analysis, the extent to which a firm uses an idea to push forward its productivity depends on the propinquity of the idea to the firm's technology class. A firm may wish to sell an idea that is not close to its own class. It can do so by using a patent agent. Analogously, the firm might want to purchase an idea through a patent agent if it fails to innovate. Due to search frictions, it may take time for a patent agent to find a buyer for a patent. Also, a patent may not be the perfect match for a buyer. R&D by firms leads to growth in the model. Additionally, there is a spillover effect from ideas. A balanced growth path for the model is explicitly characterized. A unique invariant firm-size distribution exists despite the fact that the distribution for productivity across firms is continually fanning out.

The model is calibrated in Section 4 so that it matches certain features of the U.S. aggregate economy, such as the average rate of growth, the long-run interest rate, the share of R&D in GDP, etc. It is also fit to match some facts, presented in Section 3, from the micro data on patents for U.S. public firms. Three such facts are descriptive: the share of patents that are sold, the empirical duration distribution, and the reduction in distance between a patent and its owner's line of business following a sale. Additionally, some facts from panel-data regression analysis are assembled and targeted using an indirect inference strategy. First, it is shown that a firm's market value is positively related to its citation-adjusted stock of patents, controlling for size and other things. Presumably, patents are valuable because they protect important ideas. Second, the closeness of the patents in a firm's portfolio to the firm's line of business matters for market value. Third, the more distant a patent is to a firm's line of business, the more likely it is to be sold. Fourth, a patent is closer to the buyer's line of business than to the seller's.

Clearly, a market for patents affects the incentive to do R&D. On the one hand, the fact that an idea, which is not so useful for the innovator's own production, can be sold raises the return from engaging in R&D. On the other hand, the fact that a firm can buy an idea reduces the reward from doing R&D. A goal of the analysis is to examine how a patent market affects R&D and, hence, growth. This is done in Section 5.

To gauge the importance of the patent market for economic growth and welfare, a sequence of structured thought experiments is undertaken in Section 5.

First, the rate of contact between buyers and sellers in the market is reduced to zero, which is equivalent to shutting down the market. In the absence of the patent market, the equilibrium steady-state growth rate goes from its benchmark value of 2.08% down to 2.02%, resulting in a welfare reduction of 1.18% in consumption equivalent terms. Next, the efficiency of the patent market is successively increased. It is shown that a faster rate of contact between buyers and sellers, where a buyer can find a seller without any delay, increases the growth rate up to 2.46% and leads to a welfare gain of 5.97% relative to the benchmark economy (measured in terms of consumption). In addition, if each seller is matched with the perfect buyer for their patent, then the growth rate increases to 3.05% and a welfare improvement of 14.3% materializes. Last, if the ideas that firms produce are perfectly suited for their own production process (this corresponds to a situation where there is no mismatch between a firm and the idea that it generates), then the growth rate is 3.38%, which results in a welfare gain of 17.8% compared with the baseline model. So, efficiency in the market for patents matters.

Two concerns arise with the focus on patents. First, ideas may be transferred via other means, in particular, licensing. The empirical analysis conducted in Section 3 controls for this, to the extent possible. Additionally, the model simulation is redone in Section 6 to allow for ideas to be transferred through licensing as well as patents. The results are not affected in a significant way. Second, perhaps some patents are bought and sold for reasons surrounding litigation. Such sales may have little to do with the transfer of knowledge or increasing productivity. A firm may buy an intrinsically worthless patent to fend off potential litigation, or perhaps to earn profits by threatening litigation (patent trolls). The empirical analysis in Section 3 also attempts to control for this. Additionally, as a robustness check, the model is re-simulated in Section 6 using data from low-litigation sectors. Again, the results appear to be immune to this.

The market for patents is often thought of as being inefficient and illiquid. Buying and selling intellectual property is a difficult activity. Each patent is unique. It may not be readily apparent who the potential buyers and competing sellers even are, especially in situations where enterprises desire to keep their business strategies secret. Buyers and sellers may have very different valuations about the worth of a patent. Patents are often sold through intermediaries. This motivates the search-theoretic framework presented here.

Historically, patent agents were often lawyers. Dealing with both patent buyers and sellers, they understood both sides of the market. Inventors used them to file patent applications. So, the lawyers became acquainted with the new technologies that were around. Buyers used them to vet the merits of new technologies. Hence, the lawyers were familiar with the types of patents that were likely to be marketable. This led naturally to lawyers acting as intermediaries in patent sales. Edward Van Winkle typifies the business. He was a patent agent at the beginning of the 20th century. Van Winkle was a mechanical engineer who acquired a law degree by correspondence course. He was well suited to

provide advice on the legal and technical merits of inventions for his clients on both sides of the market. Van Winkle cultivated a network of businessmen, inventors, and other lawyers. [Lamoreaux and Sokoloff \(2003\)](#) detailed how he brokered various types of deals with the buyers and sellers of patents. They also documented for the period 1870 to 1910 an increased tendency for inventors (especially the more productive ones) to use specialized registered patent agents to handle transactions associated with their patents.

While today's market for patents is sizeable, it can be regarded as being thin due to the specialized nature of the knowledge that is embodied in each patent. Thus, the patent market is highly specialized. To date, online intellectual property platforms have failed to arbitrage the market. The sensitivity of intellectual property makes potential buyers and sellers reluctant to reveal information online; they prefer face-to-face dealings with the other party. Also, some buyers may perceive a lemons problem: if the patents were truly valuable, then the sellers should be able to profit by developing the idea themselves or by selling it directly to interested parties.

1.2. *Relationship to the Literature*

How does the current paper relate to the literature? This is discussed now. On the theory side, the model developed here is in a class of its own, but like all work it is inspired by some important predecessors. The paper contributes to the endogenous growth literature. Ever since [Romer's \(1986\)](#) classic paper, economists have been concerned with how knowledge affects economic growth. The cue for a spillover effect from ideas is in the [Romer \(1986\)](#) growth model. The notion that a firm can push forward its productivity by incorporating new ideas in its production process is in [Aghion and Howitt \(1992\)](#). Unlike [Aghion and Howitt \(1992\)](#), this is done here in a competitive environment.

Recent attention has been directed to developing the micro-foundations of how new ideas spread in an economy. Some work stresses technology diffusion via innovation and imitation (e.g., [Jovanovic and MacDonald \(1994\)](#), and [König, Lorenz, and Zilibotti \(2012\)](#)). Other research emphasizes matching and other frictions in the transfer of ideas. (See, for instance, [Benhabib, Perla, and Tonetti \(2014\)](#), [Chiu, Meh, and Wright \(2011\)](#), [Lucas and Moll \(2014\)](#), and [Perla and Tonetti \(2014\)](#).) The work here emphasizes matching frictions. It differs from the above papers in a number of significant ways. First, the focus is on an economy where growth is driven by *heterogeneous* ideas that are invented by firms. A firm may not be able to make the best use of the idea it discovers. Second, firms can trade their ideas in a *market* subject to matching frictions. Third, while the growth literature has mainly been theoretical, the current research uses micro data on *patent reassignments* to motivate and discipline the analysis.⁵

⁵Perhaps the closest theoretical work to the current research is by [Chiu, Meh, and Wright \(2011\)](#). Ideas are *homogeneous* in their framework, so there cannot be any misallocation. They

The present paper highlights the importance of complementarity (as measured by distance) between the existing knowledge stock of the firm and new patents. These findings naturally relate to work on diversification. In a classic study on diversification and integration, Gort (1962, p. 108) stated “when faced with a choice among activities that would be equally attractive if they were technologically equidistant from the primary one, a firm will usually undertake those for which technical propinquity to the primary activity is greatest.” Gort (1962) provided some early evidence in support of this hypothesis. Figueroa and Serrano (2013) examined the empirical significance of this idea for patenting and licensing activities.

On the empirical side, the data employed here were first used by Serrano (2010, 2015). He used the fraction of self-citations as a proxy for the fit of an idea to an inventing firm and documented that patents that are not a good fit are more likely to be sold on the market by the inventing firm. A new metric for measuring the distance between ideas and firms is proposed here. Serrano’s (2010, 2015) findings are confirmed. Additionally, new facts on the relationship between a firm’s market value and its distance-adjusted patent portfolio are presented. Also, it is shown how the distance between an idea and its owner changes upon sale. The micro data facts that are obtained from the U.S. data are then used here to discipline a search-based endogenous growth model. The model is employed to quantify the misallocation of ideas in the U.S. economy and the contribution of the patent market to economic growth.

The focus on mismatch in ideas connects with recent work on misallocation (see, for instance, Acemoglu, Akcigit, Bloom, and Kerr (2013), Guner, Ventura, and Xu (2008), Hsieh and Klenow (2009), and Restuccia and Rogerson (2008)). That literature has mainly focused on factor misallocations, particularly the allocation of capital and labor across establishments. The current work complements this literature by focusing on differences in total factor productivity that may arise due to a misallocation of ideas, which are a direct ingredient in productivity. Ideas are not necessarily born to their best users. The existence of a market for ideas and its efficiency can have a major impact on mitigating any initial misallocation. Thus, the presence of a market for ideas may contribute significantly to productivity growth. Addressing this question is the focus of the current paper.

2. MODEL

The theoretical model with perfectly competitive firms is introduced now. The goal is to focus on the potential misallocation of ideas and its conse-

are produced by inventors who cannot commercialize them, so *all* ideas are sold. Firms cannot do R&D, hence they must purchase an idea to produce. There are search frictions in their setup: an inventor must find a entrepreneur in order to sell his idea. Their work emphasizes financial frictions. In particular, an entrepreneur must have cash on hand to buy an idea. Last, no empirical or quantitative work is done.

quences for growth and welfare; therefore, the model abstracts from monopoly distortions. Another interesting feature of this setting is that patents serve a new role in this economy: the possibility for trading ideas. Some ideas are better than others for a firm. In the analysis, there are two types of ideas: to wit, *d*-type and *n*-type. The worth of a *d*-type idea depends on the distance of the idea to a firm's main line of business. The closer the idea, the more valuable it is. The worth of an *n*-type idea is unrelated to the distance between the idea and the firm's line of business. To obtain a *d*-type idea, a firm must invest resources, either through R&D or by buying a patent on the market. By contrast, a firm may discover an *n*-type idea through serendipity for free. The productivity of both types of ideas depends upon the general pool of knowledge in the economy; that is, through osmosis, some components of ideas become part of the ether in technology space.⁶

2.1. *Environment*

Consider an economy, where time flows discretely, with a continuum of firms of unit measure. The firms produce a homogeneous final good using capital and labor. Each firm belongs permanently to some technology class j that resides on a circle with radius $1/\pi$. At each point on the technology circle, there are firms of density $1/2$. A firm enters the period with a level of productivity z . At the beginning of a period, each firm develops a *d*-type idea with an *endogenous* probability i . The *d*-type innovation will be patented and belongs to some technology class k on the circle. The distance between the firm's own technology class, j , and the innovation, k , is denoted by $d(j, k)$. This represents the length of the shortest arc between j and k . Transform this distance measure into a measure of technological propinquity, $x = 1 - d(j, k)$, defined on $[0, 1]$. A high value for x indicates that the innovation is close to the firm's technology class. The firm will keep or sell the *d*-type patent depending on the value for x . The higher x is, the bigger will be the boost to the firm's productivity, if the firm decides to keep the idea. The value of x is drawn from the distribution function $X(x)$. The technology circle is illustrated in Figure 2. Just before production begins, an *n*-type idea arrives with an *exogenous* probability p . The worth of an *n*-type idea is unrelated to a firm's technology class. The analysis will focus on a symmetric equilibrium around the technology circle. In a symmetric equilibrium, at each point on the circle the distribution of firms is the same. Analyzing one point on the circle is the same as analyzing any other, so there is no need to carry around a location index.

⁶A simplified version of the model that connects in a more elementary manner the efficiency of the patent market, and the propinquity of an idea with the firm's line of business, to economic growth was put forward in [Akcigit, Celik, and Greenwood \(2015, Appendix 12\)](#).




FIGURE 2.—The technology circle (left panel) and the timing of events (right panel) for d -type ideas. Note that n -type ideas arrive after the market for d -type patents closes.

Firms produce output, o , at the end of a period according to the production process

$$(1) \quad o = (e' z')^\zeta k^\kappa l^\lambda, \quad \text{with} \quad \zeta + \kappa + \lambda = 1,$$

where k and l are the amounts of capital and labor used in production and z' is its end-of-period productivity. The variable e' is a firm-specific idiosyncratic production shock. It is drawn at the end of each period from a log-normal distribution with $E[e'] = 1$ and a standard deviation represented by $\text{STD}(\ln e')$.⁷ Labor is hired at the wage rate w . There is one unit of labor available in the economy. Capital is hired at the rental rate \tilde{r} . Observe that there are diminishing returns in capital and labor. Hence, there are profits from producing. These rents are increasing in the firm's productivity, z' . This provides an incentive to do R&D to improve z' . The exponent ζ on $e' z'$ is an innocuous normalization that results in profits being linear in $e' z'$, as is shown below.

A firm's end-of-period productivity, z' , evolves according to the law of motion

$$(2) \quad z' = L(z, x, b; \mathbf{z}) = z + \gamma_d x z + \gamma_n b z.$$

Here z is the firm's initial productivity level. The second term gives the increment to productivity from obtaining a d -type patent, where x is the technological propinquity of the patent to the firm and \mathbf{z} is mean of the productivity distribution in the economy at the beginning of the period. The closer a d -type innovation is to a firm's own technology class, as represented by a larger x , the bigger will be the increase in productivity, $\gamma_d x z$. The third term gives the gain in productivity from acquiring an n -type idea, where $b \in \{0, 1\}$. The expected

⁷The e' shock implies that employment, l , will not be a perfect predictor of a firm's market value. This property is important for the quantitative analysis and is discussed in Section 4.

value of b is given by $E[b] = p$. Once an idea is blended into a firm's production process, within the firm's permanent technology class, it loses its individual identity. This assumption implies that there is no need to keep track of a firm's portfolio of patents, which would vary by technology class and age; doing so would be an insurmountable task.

The higher is the economy-wide baseline level of productivity, z , the more valuable a patent is for increasing productivity. This is true for both d -type and n -type patents. Note that z introduces a knowledge externality in this economy. Although not modeled formally, this could be because either some forms of knowledge can only be imperfectly protected or because the patents protecting them have expired so that the knowledge formerly embodied in the patents is now freely available for all. Since n -type patents arrive with exogenous probability p , the firm will benefit from spillovers in a probabilistic sense, even if the firm does not invest any resources in R&D. This is not true for d -type patents, as will be seen. Later, the notation $L(z, 0, b; z)$ will be used to signify the situation where the firm's productivity is not incremented by a d -type innovation in the current period, which is equivalent to setting $x = 0$. One might think that a firm would try to discover a d -type idea that is close to its line of business. As was mentioned, the propinquity of a d -type idea to the firm, x , is drawn from the distribution $X(x)$. In the quantitative analysis, this is taken to be the empirical distribution. Hence, the propinquity of ideas to their inventors will be the same as in the data. It turns out that z is also the aggregate state variable in this economy, a fact shown later. That is, only the mean of the distribution for the z 's across firms and the evolution of this mean over time matter for the analysis. Assume that z evolves according to the deterministic aggregate law of motion

$$(3) \quad z' = T(z).$$

Now, at the beginning of a period, firms pick the probability of discovering a d -type idea, i . They do this according to the convex cost function

$$(4) \quad C(i; z) = \chi z^{\xi/(\xi+\lambda)} i^{1+\rho} / (1 + \rho).$$

Cost rises in lock-step fashion with average productivity, z , in the economy. It will be established later that wages, w , are proportional to z and grow along a balanced growth path at the same rate as $z^{\xi/(\xi+\lambda)}$. As will be seen, this ensures that, along a balanced growth path, the ratio of aggregate R&D expenditures to GDP remains constant. Aggregate productivity will be a function of the aggregate state of the world represented by z . A firm that successfully innovates can either keep or sell its idea to a patent agent. A firm that does not innovate can try to buy a patent from an agent. A patent on the market survives over time with probability σ . In the analysis, σ will be set so that patents have the same expected life as in the U.S. data. But, by letting a patent die stochastically

in this fashion, instead of deterministically, there is no need to keep track of a patent's age, a huge simplification.

d -type ideas can be bought and sold on a patent market. A firm that fails to come up with a d -type idea can try to buy one from a patent agent. Innovators are not allowed to buy d -type patents. A firm that draws a d -type idea may sell the associated patent to a patent agent at the price q . This price is determined on a competitive market. Once a patent is sold to an agent, the seller cannot use it in the future. A patent agent can only handle one d -type idea at a time. The introduction of patent agents simplifies the analysis. Without this construct, the analysis would have to keep track of the portfolio of patents that each firm has for sale. This technical construct is imposed without apology, as in the real world many patents are sold through agents, as was discussed.

Let n_a and n_b represent the numbers of agents and buyers in the market for d -type patents. The total number of matches in the market is given by the matching function

$$M(n_a, n_b) = \eta n_a^\mu n_b^{1-\mu}.$$

The matches are completely random. Thus, the odds that an agent will find a buyer are given by

$$m_a\left(\frac{n_a}{n_b}\right) = \frac{M(n_a, n_b)}{n_a} = \eta \left(\frac{n_b}{n_a}\right)^{1-\mu},$$

and similarly that a buyer will find an agent by

$$m_b\left(\frac{n_a}{n_b}\right) = \frac{M(n_a, n_b)}{n_b} = \eta \left(\frac{n_a}{n_b}\right)^\mu.$$

This search friction could reflect many things: the hardship of matching buyers and sellers in a thin market for a complicated product or the difficulty of a buyer assessing the quality of a patent for his line of business, inter alia.

The ratio of potential sellers to buyers, n_a/n_b , reflects the slackness of the market. Since agents and buyers are matched randomly, the propinquity between the buyer's technology class and the class of the d -type patent being sold is a random variable. A buyer will incorporate a d -type patent that he purchases into his production process in accordance with the above law of motion for z . The price of the d -type patent is determined by Nash bargaining between the agent and buyer. Represent this price by $p = P(z, x; \mathbf{z})$. The negotiated price will depend on the propinquity of the patent, x , and the state of the buyer's technology, z . The bargaining power of the agent is given by ω . In contrast, the price at which a firm sells its d -type patent to an agent is fixed at q , because the agent does not know who he will sell the patent to in the future. The timing of events in the market for d -type patents is portrayed in the right panel of Figure 2. Last, after the d -type patent market closes, an n -type idea

may arrive to a firm. For the moment, assume that n -type ideas are not traded. A market for n -type ideas is appended onto the model in Section 3.4.

2.2. The Representative Consumer/Worker

In the background of the analysis is a representative consumer/worker. This individual supplies one unit of labor inelastically. The person owns all of the firms in the economy. He also rents out the capital used by firms. Thus, he will earn income from wages, profits, and rentals. Capital depreciates at the rate δ . The real return earned by renting capital is $1/r$. (That is, r is the reciprocal of the gross interest rate. It will play the role of the discount factor in the Bellman equations formulated below.) The individual is assumed to have a momentary utility function of the form $U(c) = c^{1-\varepsilon}/(1-\varepsilon)$, where c is his consumption in the current period and ε is the coefficient of relative risk aversion. He discounts the future at rate β . Last, the representative consumer/worker's goal in life is to maximize his discounted lifetime utility. Since this problem is entirely standard, it is not presented.

2.3. Firms: Buyers, Keepers, and Sellers

A firm hires labor, l , at the wage rate, w , and capital, k , at the rental rate, $\tilde{r} \equiv 1/r - 1 + \delta$, to maximize profits. It does this at the end of each period after seeing the realized values for e' and z' . Thus, its decision problem is

$$e'\Pi(z'; \mathbf{z}) = \max_{k,l} [(e'z')^\zeta k^\kappa l^\lambda - \tilde{r}k - wl],$$

where $e'\Pi(z'; \mathbf{z})$ is the profit function associated with the maximization problem; the fact that this function is multiplicative in e' is established momentarily. The first-order conditions to this maximization problem imply that

$$(5) \quad k = \kappa \frac{o}{\tilde{r}},$$

and

$$(6) \quad l = \lambda \frac{o}{w}.$$

Using (1), (5), and (6), it follows that profits are given by

$$(7) \quad e'\Pi(z'; \mathbf{z}) = (1 - \kappa - \lambda)o = e'z'(1 - \kappa - \lambda) \left[\left(\frac{\kappa}{\tilde{r}} \right)^\kappa \left(\frac{\lambda}{w} \right)^\lambda \right]^{1/\zeta}.$$

Again, in equilibrium the rental and wage rates, \tilde{r} and w , will be functions of the aggregate state of the world, \mathbf{z} . Note that profits are increasing in z' when there are diminishing returns to scale ($1 - \kappa - \lambda < 1$). This provides an incentive to innovate.

The value function for a firm that desires to buy a patent will now be formulated. To this end, let $V(z; \mathbf{z})$ represent the expected present value of a firm that currently has productivity z and is about to learn whether or not it has come up with a d -type idea. Due to the focus on symmetric equilibrium, there is no need ever to record the firm's location on the technology circle. Now, suppose that the firm does not innovate. Then, it will try to buy a d -type patent. With probability $1 - m_b(n_a/n_b)$, it will fail to find a patent agent. In that circumstance, the firm's productivity will remain at z ; this is equivalent to setting $x = 0$ in (2). It may still acquire an n -type patent before the start of production, though, which would allow productivity to increase by $\gamma_n \mathbf{z}$. The expected value of the firm, before the n -type patent shock, is $E[\Pi(L(z, 0, b; \mathbf{z}); \mathbf{z})] + rE[V(L(z, 0, b; \mathbf{z}); \mathbf{z}')]—$ recall that $b \in \{0, 1\}$ is a random variable connected with the n -type idea that takes the value 1 with probability p and that $E[e'] = 1$.

With probability $m_b(n_a/n_b)$, the firm will meet an agent selling a d -type patent with propinquity x . Two things can happen here: either the firm buys a d -type patent from the agent or it does not. The d -type patent sells at the price $p = P(z, x; \mathbf{z})$, which is a function of the buyer's type, z , as well as the propinquity of the patent to the firm's technology class, x . The determination of the d -type patent price is discussed below. The firm will only buy the patent if it yields a higher payoff than what it will obtain if it does not buy it. If the firm buys a patent, its productivity will rise to $z + \gamma_d x \mathbf{z}$. Again, before production begins, the firm may also obtain an n -type patent, which would result in a further increase in productivity. The firm's expected value (before the n -type patent shock) will then move up to $E[\Pi(L(z, x, b; \mathbf{z}); \mathbf{z})] - P(z, x; \mathbf{z}) + rE[V(L(z, x, b; \mathbf{z}); \mathbf{z}')]$. If it does not buy a d -type patent, then its productivity will remain at z . The expected value of the firm will then be $E[\Pi(L(z, 0, b; \mathbf{z}); \mathbf{z})] + rE[V(L(z, 0, b; \mathbf{z}); \mathbf{z}')]$. Denote the distribution over propinquity for buying a patent from a patent agent by $D(x)$.

The expected discounted present value of the buyer, $B(z; \mathbf{z})$, is easily seen to be

$$\begin{aligned}
 (8) \quad B(z; \mathbf{z}) = & m_b \left(\frac{n_a}{n_b} \right) \int \left\{ I_a(z, x; \mathbf{z}) \{ E[\Pi(L(z, x, b; \mathbf{z}); \mathbf{z})] \right. \\
 & - P(z, x; \mathbf{z}) + rE[V(L(z, x, b; \mathbf{z}); \mathbf{z}')] \} \\
 & + [1 - I_a(z, x; \mathbf{z})] \{ E[\Pi(L(z, 0, b; \mathbf{z}); \mathbf{z})] \\
 & + rE[V(L(z, 0, b; \mathbf{z}); \mathbf{z}')] \} \} dD(x) \\
 & + \left[1 - m_b \left(\frac{n_a}{n_b} \right) \right] \{ E[\Pi(L(z, 0, b; \mathbf{z}); \mathbf{z})] \\
 & + rE[V(L(z, 0, b; \mathbf{z}); \mathbf{z}')] \},
 \end{aligned}$$

where \mathbf{z} evolves according to (3) and

$$(9) \quad I_a(z, x; \mathbf{z}) = \begin{cases} 1 \text{ (sale),} & \text{if the buyer purchases a patent,} \\ 0 \text{ (no sale),} & \text{otherwise.} \end{cases}$$

The indicator function $I_a(z, x; \mathbf{z})$, defined above, specifies whether or not the non-innovating firm will buy a d -type patent. The determination of this function is discussed below.

Turn now to the situation where the firm successfully innovates. If it decides to keep the d -type patent, then the firm's productivity will be $z + \gamma_d x \mathbf{z}$ as in (2). Productivity may still increase if the firm draws an n -type idea. Before the realization of the n -type patent shock, the firm will have the expected value $K(z + \gamma_d x \mathbf{z}; \mathbf{z})$, as given by

$$(10) \quad K(z + \gamma_d x \mathbf{z}; \mathbf{z}) = E[\Pi(L(z, x, b; \mathbf{z}); \mathbf{z})] + rE[V(L(z, x, b; \mathbf{z}); \mathbf{z}')],$$

where again \mathbf{z} evolves according to (3) and $b \in \{0, 1\}$ is a random variable. Alternatively, it can sell the d -type patent to an agent. Then, its productivity will remain at z (unless it subsequently draws an n -type idea). The value of a seller, $S(z; \mathbf{z})$, is

$$(11) \quad S(z; \mathbf{z}) = E[\Pi(L(z, 0, b; \mathbf{z}); \mathbf{z})] + \sigma q + rE[V(L(z, 0, b; \mathbf{z}); \mathbf{z}')].$$

Once the seller puts a d -type patent up for sale at the beginning of the period, it expires with probability $1 - \sigma$. A firm that innovates will either keep or sell its d -type patent depending on which option yields the highest value. Given this, it is easy to see that the decision to keep or to sell a patent can be formulated as

$$(12) \quad I_k(z, x; \mathbf{z}) = \begin{cases} 1 \text{ (keep),} & \text{if } K(z + \gamma_d x \mathbf{z}; \mathbf{z}) > S(z; \mathbf{z}), \\ 0 \text{ (sell),} & \text{otherwise.} \end{cases}$$

2.3.1. The Decision to Innovate

The firm's decision to innovate is now cast. With probability i the firm discovers a d -type idea and with probability $1 - i$ it does not. The firm chooses the probability of discovering a d -type idea subject to the convex cost function $C(i; \mathbf{z})$. Hence, write the innovation decision as

$$(13) \quad V(z; \mathbf{z}) = \max_i \left\{ i \int \{ I_k(z, x; \mathbf{z}) K(z + \gamma_d x \mathbf{z}; \mathbf{z}) + [1 - I_k(z, x; \mathbf{z})] S(z; \mathbf{z}) \} dX(x) + (1 - i) B(z; \mathbf{z}) - C(i; \mathbf{z}) \right\}.$$

The first-order condition associated with this problem is

$$\begin{aligned} & \int \{I_k(z, x; \mathbf{z})K(z + \gamma_d x \mathbf{z}; \mathbf{z}) \\ & \quad + [1 - I_k(z, x; \mathbf{z})]S(z; \mathbf{z})\} dX(x) - B(z; \mathbf{z}) \\ & = C_1(i; \mathbf{z}) \end{aligned}$$

(where C_1 is the derivative of C with respect to i), so that

$$\begin{aligned} (14) \quad i &= R(z; \mathbf{z}) \\ &= C_1^{-1} \left(\int \{I_k(z, x; \mathbf{z})K(z + \gamma_d x \mathbf{z}; \mathbf{z}) \right. \\ & \quad \left. + [1 - I_k(z, x; \mathbf{z})]S(z; \mathbf{z})\} dX(x) - B(z; \mathbf{z}; \mathbf{z}) \right). \end{aligned}$$

2.4. Patent Agents

Turn now to the problem of a patent agent. It buys a d -type idea at the competitively determined price q . With probability $m_a(n_a/n_b)$ it will meet a potential buyer on the market and with probability $1 - m_a(n_a/n_b)$ it will not. Denote the distribution of buyers by $G(z)$. The value for an agent, A , with a patent is thus given by

$$\begin{aligned} (15) \quad A(\mathbf{z}) &= m_a \left(\frac{n_a}{n_b} \right) \int \int \{I_a(z, x; \mathbf{z})P(z, x; \mathbf{z}) \\ & \quad + [1 - I_a(z, x; \mathbf{z})]r\sigma A(\mathbf{z}')\} dG(z) dD(x) \\ & \quad + \left[1 - m_a \left(\frac{n_a}{n_b} \right) \right] r\sigma A(\mathbf{z}'), \end{aligned}$$

where $I_a(z, x; \mathbf{z})$ is specified by (9) and is defined formally shortly below. The price of a d -type patent is determined via Nash bargaining. Specifically, p is determined in accordance with

$$\begin{aligned} & \max_p \{E[\Pi(L(z, x, b; \mathbf{z}); \mathbf{z})] - p + rE[V(L(z, x, b; \mathbf{z}); \mathbf{z}')]\} \\ & \quad - E[\Pi(L(z, 0, b; \mathbf{z}); \mathbf{z})] - rE[V(L(z, 0, b; \mathbf{z}); \mathbf{z}')]\}^{1-\omega} \\ & \quad \times [p - r\sigma A(\mathbf{z}')]^\omega. \end{aligned}$$

The first term in braces gives the buyer's surplus. This gives the difference between the value of the firm when it secures a d -type patent and the value when

it does not. The second term details the seller's surplus. In standard fashion,

$$(16) \quad p = P(z, x; \mathbf{z}) \\ = \omega \{ E[\Pi(L(z, x, b; \mathbf{z}); \mathbf{z})] + rE[V(L(z, x, b; \mathbf{z}); \mathbf{z}')] \\ - E[\Pi(L(z, 0, b; \mathbf{z}); \mathbf{z})] - rE[V(L(z, 0, b; \mathbf{z}); \mathbf{z}')] \} \\ + (1 - \omega)r\sigma A(\mathbf{z}'),$$

whenever both the buyer's and seller's surpluses are positive. The price lies between $r\sigma A(\mathbf{z}')$ and $E[\Pi(L(z, x, b; \mathbf{z}); \mathbf{z})] + rE[V(L(z, x, b; \mathbf{z}); \mathbf{z}')] - E[\Pi(L(z, 0, b; \mathbf{z}); \mathbf{z})] - rE[V(L(z, 0, b; \mathbf{z}); \mathbf{z}')]$; if the former is above the latter, then no solution exists. Now, define $I_a(z, x; \mathbf{z})$ in the following manner:

$$(17) \quad I_a(z, x; \mathbf{z}) \\ = \begin{cases} 1, & \text{if } r\sigma A(\mathbf{z}') \leq p \leq E[\Pi(L(z, x, b; \mathbf{z}); \mathbf{z})] \\ & + rE[V(L(z, x, b; \mathbf{z}); \mathbf{z}')] - E[\Pi(L(z, 0, b; \mathbf{z}); \mathbf{z})] \\ & - rE[V(L(z, 0, b; \mathbf{z}); \mathbf{z}')], \\ 0, & \text{otherwise.} \end{cases}$$

2.5. Symmetric Equilibrium Along a Balanced Growth Path

The focus of the analysis is solely on a symmetric equilibrium along a balanced growth path. A formal analysis of the model's balanced growth path is contained in Theory Appendix A of the Supplemental Material. Before starting, define the aggregate level of productivity, \mathbf{z} , its gross rate of growth, \mathbf{g} , and the aggregate level of innovation, \mathbf{i} , by

$$(18) \quad \mathbf{z} \equiv \int z dZ(z), \quad \mathbf{g} \equiv \frac{\int z' dZ'(z')}{\int z dZ(z)}, \quad \text{and} \quad \mathbf{i} \equiv \int R(z; \mathbf{z}) dZ(z).$$

In equilibrium, the demand for labor must equal the supply of labor. Recall that there is one unit of labor in the economy. Let $Z'(z')$ represent the end-of-period distribution of z' across firms. Now, using (1), (5), and (6), it is easy to deduce that the labor, l , demanded by a firm is given by

$$(19) \quad l = \left(\frac{\kappa}{\tilde{r}} \right)^{\kappa/\zeta} \left(\frac{\lambda}{w} \right)^{(\zeta+\lambda)/\zeta} e' z'.$$

Equilibrium in the labor market then implies that

$$\int \left(\frac{\kappa}{\tilde{r}} \right)^{\kappa/\zeta} \left(\frac{\lambda}{w} \right)^{(\zeta+\lambda)/\zeta} z' dZ'(z') = 1,$$

where the fact that $E[e'] = 1$ has been used. This implies that the aggregate wage rate, w , is given by

$$(20) \quad w = \lambda \left(\frac{\kappa}{\tilde{r}} \right)^{\kappa/(\xi+\lambda)} \left[\int z' dZ'(z') \right]^{\xi/(\xi+\lambda)} = \lambda \left(\frac{\kappa}{\tilde{r}} \right)^{\kappa/(\xi+\lambda)} \mathbf{z}'^{\xi/(\xi+\lambda)}.$$

The wage rate, w , depends on the mean of the end-of-period productivity distribution across firms, $\mathbf{z}' \equiv \int z' dZ'(z')$.

Next, suppose that there is free entry by agents into the market for d -type patents. This dictates that the price q will be determined by

$$(21) \quad q = A(\mathbf{z}).$$

To complete the description of a symmetric balanced growth equilibrium, the distribution over propinquity for patent agents, or $D(x)$, must be specified. It is uniform in a symmetric equilibrium. Recall that a firm's permanent location in the technology space is represented by a point on the circle. Think about a buyer located at the top of the circle. Suppose that a set of firms on some tiny arc jk to the left of top are selling patents of mass λ that are of distance between 0 and ε away from the top. Now take any other arc lm of equal length even further to the left of top. The start of this second arc has distance $d(j, l)$ from the start of the first one. In a symmetric equilibrium, there will be on the second arc, for all practical purposes, an identical set of firms selling patents of mass λ that are of distance between $d(j, l)$ and $d(j, l) + \varepsilon$ away from the top.

2.5.1. Some Features of a Balanced Growth Path

Along a balanced growth path, consumption, investment, output, profits, wages, and the selling and buying prices for d -type patents will all grow at a constant rate. Also, the interest factor and rental rate on capital are constant. Assuming that this is the case, then it is easy to deduce from (20) that wages must grow at the gross rate $\mathbf{g}^{\xi/(\xi+\lambda)}$. Aggregate output and profits will grow at this rate, too, as can be inferred from (7). Given the assumption that tastes are isoelastic, the interest factor and rental rate on capital are given in standard fashion by

$$(22) \quad r = \beta / \mathbf{g}^{\varepsilon \xi / (\xi + \lambda)},$$

and

$$(23) \quad \tilde{r} = \mathbf{g}^{\varepsilon \xi / (\xi + \lambda)} / \beta - 1 + \delta,$$

where again ε is the coefficient of relative risk aversion. By substituting the solution for wages, as given by (20), into the demand for labor, (19), it can be seen that a firm's employment is proportional to z'/\mathbf{z}' . Since, on average, one would expect that z' will be growing at the same rate as \mathbf{z}' , this suggests that a

stationary firm-size distribution exists. (For a formal proof, see Akcigit, Celik, and Greenwood (2015, Appendix 11).)

It turns out that along a balanced growth path, the indicator functions $I_k(z, x; \mathbf{z})$ and $I_a(z, x; \mathbf{z})$ can be represented by simple threshold rules for x that do not depend on either z or \mathbf{z} . In particular,

$$I_k(z, x; \mathbf{z}) = \begin{cases} 1 \text{ (keep)}, & x > x_k, \\ 0 \text{ (sell)}, & \text{otherwise,} \end{cases} \quad \text{and}$$

$$I_a(z, x; \mathbf{z}) = \begin{cases} 1 \text{ (sale)}, & x > x_a, \\ 0 \text{ (no sale)}, & \text{otherwise.} \end{cases}$$

That is, an innovating firm keeps its d -type idea when $x > x_k$ and sells otherwise. Analogously, a sale between a buyer and a patent agent occurs if and only if $x > x_a$.

3. EMPIRICAL ANALYSIS

3.1. Data Sources

This section details data sources and variable constructions. For further information, please see Empirical Appendix B of the Supplemental Material.

NBER-USPTO Utility Patents Grant Data (PDP). The core of the empirical analysis draws from the NBER-USPTO Patent Data Project (PDP). Patents are exclusionary rights, granted by national patent offices, to protect a patent holder for a certain amount of time, conditional on sharing the details of the invention. The PDP data contain detailed information on 3,210,361 utility patents granted by the U.S. Patent and Trademark Office between the years 1976 and 2006. A patent has to cite another patent when the former has content related to the latter. When patent A cites patent B, this particular citation becomes both a *backward* citation made by A to B and a *forward* citation received by B from A. Moreover, the PDP contains an International Patent Classification (IPC) code for each patent that helps identify where it lies in the technology space.⁸ Extensive use of the forward and backward citations is made, as well as the IPC codes assigned to each patent, to determine a patent's location in the technology space, its distance to a firm's location in the technology spectrum, and also to proxy for a patent's quality. The exact methodology followed to construct these measures is detailed below.

Patent Reassignment Data (PRD). The second source of data comes from the recently released USPTO patent assignment files retrieved from Google

⁸The USPTO originally assigns each patent to a particular U.S. Patent Classification (USPC), which is a system used by the USPTO to organize all patents according to their common technological relevances. The PDP also assigns an IPC code to each patent using the original USPC and a USPC-IPC concordance based on the International Patent Classification Eighth Edition.

Patents Beta. (This data is now maintained by the USPTO.) This data set provides detailed information on the changes in patent ownership for the years 1980 to 2011. The records include 966,427 patent reassignments not only due to *sales*, but also due to *mergers*, *license grants*, *splits*, *mortgages*, *collaterals*, *conversions*, *internal transfers*, etc. Reassignment records are classified according to a search algorithm that looks for keywords, such as “assignment,” “purchase,” “sale,” and “merger,” and assigns them to their respective categories. Through this process, 99% of the transaction records are classified into their respective groups—see Empirical Appendix B for more information.

Compustat North American Fundamentals (Annual). In order to assess the impact of patents and their technological distance on firm moments, such as stock market valuation, the PDP patent data are linked to Compustat firms. The focus is on the balance sheets of Compustat firms between the years 1974 and 2006, retrieved from Wharton Research Data Services. The Compustat database and the NBER PDP database are connected using the matching procedure provided in the PDP data.

Lex Machina Database on Patent Litigations. The information on litigated patents is obtained from Lex Machina. It is the most comprehensive database on patent litigations since 2000. Lex Machina obtains its data on a daily basis from (i) the administrative database of the United States federal courts, (ii) all United States District Courts’ websites, (iii) the International Trade Commission’s (EDIS) website, and (iv) the USPTO’s websites.

Derwent LitAlert Database on Patent Litigations. For litigation information before 2000, the Derwent LitAlert Database is used. Further description about this data set can be found in Galasso, Schankerman, and Serrano (2013).

Carnegie Mellon Survey (CMS) on Industrial R&D. The sector-level licensing information is drawn from the CMS. This data set is one of the rare R&D surveys in the United States that contains information on the licensing activities of firms. The CMS contains 1,478 randomly selected R&D labs of manufacturing firms, stratified by three-digit SIC industry codes. All labs are located in the United States. In the survey, the firms are asked to report the most important reason for applying for their product patent, where one of the answers is “to obtain revenue through licensing.” The percentage of firms picking this answer is aggregated to two-digit SIC industry classifications, which results in a sector-level licensing intensity measure. More information can be found in Cohen, Nelson, and Walsh (2000).

The empirical analysis requires the construction of a notion of distance in the technology space. For that purpose, the citation patterns across IPC technology fields are utilized. The PDP contains the full list of citations with the identity of citing and cited patents. Since the data also contain the IPC code of each patent, the percentage of outgoing citations from one technology class to another is observable. Using this information, a metric, discussed below, is constructed to gauge the distance between a new patent and a firm’s location in the technology spectrum.

In what follows, for each empirical fact the best and largest possible sample is used. For instance, for the firm value regressions, all patents that are matched to the Compustat sample are utilized. Similarly, to describe the change from seller to buyer, all patents for which the buyer and seller could be uniquely identified are used. Therefore, even though the samples vary across different empirical facts, this approach delivers the most reliable results.

3.2. *Technological Propinquity*

The notion of technological propinquity between a patent and a firm is now formalized. Think about a patent as lying within some technological class. Call this technology class X . Empirically, this can be represented by the first two digits of its International Patent Classification (IPC) code. Now, one can measure how close two patent classes, X and Y , are to each other. To do this, let $\#(X \cap Y)$ denote the number of all patents that cite patents from technology classes X and Y simultaneously. Let $\#(X \cup Y)$ denote the number of all patents that cite either technology class X and/or Y . Then, the following symmetric distance metric can be constructed:

$$d(X, Y) \equiv 1 - \frac{\#(X \cap Y)}{\#(X \cup Y)},$$

with $0 \leq d(X, Y) \leq 1$. This distance metric is intuitive. If each patent that cites X also cites Y , this metric delivers a distance of $d(X, Y) = 0$. (Also note that $d(X, X) = 0$.) If there is no patent that cites both classes, then the distance becomes $d(X, Y) = 1$. The distance between two technology classes increases, as the fraction of patents that cite both decreases. Given this metric between technology classes, a distance measure between a patent and a firm can now be constructed.

In order to measure how close a patent is to a firm in the technology spectrum, a metric needs to be devised. For this purpose, a firm's past patent portfolio is used to identify the firm's existing location in the technology space.⁹ In particular, the distance of a particular patent p to a firm f is computed by calculating the average distance of p to each patent in firm f 's patent portfolio as follows:

$$(24) \quad d_{\iota}(p, f) \equiv \left[\frac{1}{\|\mathcal{P}_f\|} \sum_{p' \in \mathcal{P}_f} d(X_p, Y_{p'})^{\iota} \right]^{1/\iota},$$

with $0 < \iota \leq 1$, and where $0 \leq d_{\iota}(p, f) \leq 1$. In this expression, \mathcal{P}_f denotes the set of all patents that were ever invented by firm f prior to patent p ,

⁹The firm's patent portfolio is defined as all inventions by the firm up to that point in time.




FIGURE 3.—Empirical distance distributions. The figure plots empirical density functions for the distance, $d_i(p, f)$, between a patent, p , and a firm's patent portfolio, f , for three values of i .

$\|\mathcal{P}_f\|$ stands for its cardinality, and $d(X_p, Y_{p'})$ measures the distance between the technology classes of patents p and p' . Note that $d(X_p, Y_{p'}) = 0$ when the firm has another patent, p' , in the same class as p . Therefore, this metric is defined only for $i > 0$. Finally, when $i = 1$, the above metric returns the average distance of p to each patent in firm f 's patent portfolio: $d_1(p, f) \equiv \|\mathcal{P}_f\|^{-1} \sum_{p' \in \mathcal{P}_f} d(X_p, Y_{p'})$, with $0 \leq d_1(p, f) \leq 1$.

The empirical distribution for this notion of distance is displayed in Figure 3 for three values of i . As can be seen, patents have heterogeneous technological distances to the inventing firms. The intermediate value, $i = 2/3$, is chosen for the subsequent analysis.¹⁰

3.3. Stylized Facts

Next, the empirical findings highlighted in the Introduction of the paper are presented. Table I provides the summary statistics. Panel A shows the summary statistics of the variables computed using Compustat firms. The distance-adjusted patent stock is constructed in a way such that each patent's contribution to the portfolio is multiplied by its distance to the firm prior to the aggregation. Specifically,

$$\sum_{p \in \mathcal{P}_f} d_i(p, f) \times \text{QUALITY}(p),$$

¹⁰The value chosen for i does not appear to make much of a difference for the analysis. For example, both the empirical and model simulation results in the paper are more or less the same when either $i = 1/3$ or $i = 1$.

TABLE I
SUMMARY STATISTICS^a

	Observations	Mean	St. Dev.
Panel A. Compustat Facts			
log market value	37,331	5.58	2.30
log employment	39,431	0.75	2.28
log patent stock	41,515	5.68	2.25
log distance-adjusted patent stock	42,269	3.38	4.23
Panel B. USPTO/NBER Patent Facts			
Patent quality	2,771,692	12.1	20.6
Patent distance	2,548,998	0.48	0.30
Litigation probability	2,772,895	0.01	0.10
Panel C. Patent Reassignment Facts			
Fraction of patents sold (at least once)	3,210,361	0.16	0.36
Number of times a patent is sold	3,210,361	0.19	0.52
Conditional duration of patent sale, yrs	421,936	5.48	4.58
Litigation and sale probability	2,772,895	0.003	0.05
Panel D. Cumulative Density			
	0 Times	1 Time	2 Times
Number of times a patent is sold	85%	97%	99%

^aPatent quality is measured by the number of patent citations corrected for truncation using the “HJT correction term” from Hall, Jaffe, and Trajtenberg (2001). “Portfolio size” is defined as the number of patents that the innovating firm has ever produced by the time of the current innovation. “Transfer duration” is measured by the grant date, with negative durations being dropped.

where $d_i(p, f)$ and $QUALITY(p)$ are the distance and quality terms for patent p . The quality of a patent is measured by the citations it has received from other patents, corrected for truncation and technology class biases using the weights offered in Hall, Jaffe, and Trajtenberg (2001).

Panel B reports the summary statistics of the USPTO/NBER patent data. As seen, the average distance between a new patent and its firm is 0.48. The so-called “garage inventors” and firms that do not have any existing patents in their portfolio are dropped when patent distance is computed. Panel C lists the summary statistics using patent reassignment data. On average, 15.6% of patents in the sample were traded at least once. The mean time to sell a patent after its grant date is 5.5 years. The average number of trades per patent is 0.2. Panel D shows that 97% of patents are traded at most one time and this number goes up to 99% when the fraction of patents that are traded at most two times are considered. Only a paltry 1.0% of patents involve litigation. The following fact summarizes this section.

FACT 1: *About 16% of patents are sold and it takes about 5.5 years to sell them, on average.*

3.3.1. Firm Market Value and Patent–Firm Distance

Are patent–firm distances important when it comes to the relationship between a firm’s patent portfolio and its value? In order to answer this question, Table II regresses “log market value” in year t on a firm’s patent portfolio, its distance-adjusted patent portfolio, and the firm’s size in the same year. The regressions also include year and firm fixed effects to rule out firm-specific properties and time trends.

As expected, column 1 shows that the patent portfolio of a firm is positively related to its stock market valuation. Presumably this is because patents are protecting knowledge that is valuable for the firm. More interestingly, a firm’s patent portfolio, once adjusted by patent distances, is negatively related to the firm’s market value. The coefficient for the distance-adjusted patent stock quantifies the loss of correlation between the patent portfolio and firm value due to the technological mismatch between the firm and its patents. In short, while the non-distance component of the patent portfolio contributes positively, the distance-related component contributes negatively to firm value. In order to interpret the results correctly, consider the ratio of the (negative) coefficient of the distance-adjusted patent stock to that of the unadjusted patent stock. The ratio of the two elasticities is 51.3%. This reflects the relative importance on market value of a shift in the distance-adjusted patent portfolio versus a change in the non-adjusted one. This ratio will be targeted in the simulation. It provides information on the importance of d -type patents relative to n -type ones.

TABLE II
FIRM MARKET-VALUE REGRESSIONS^a

	Dependent Variable: log Market Value					
	(1)	(2)	(3)	(4)	(5)	(6)
log patent stock	0.039*** (0.008)	0.039*** (0.008)	0.039*** (0.008)	0.038*** (0.009)	0.037*** (0.009)	0.071*** (0.010)
log dist-adj pat stock	−0.020*** (0.003)	−0.020*** (0.003)	−0.020*** (0.003)	−0.020*** (0.003)	−0.020*** (0.003)	−0.032*** (0.004)
log employment	0.728*** (0.008)	0.728*** (0.008)	0.736*** (0.008)	0.763*** (0.009)	0.763*** (0.009)	0.692*** (0.010)
Firm litigation intensity	no	yes	no	no	yes	no
Sector litigation intensity	no	no	yes	no	yes	no
Sector licensing intensity	no	no	no	yes	yes	no
Only renewed patents	no	no	no	no	no	yes
Observations	36,094	36,091	36,094	33,062	33,060	27,528
R^2	0.92	0.92	0.92	0.92	0.92	0.92

^aCompustat, firm-level regressions. Standard errors are in parentheses and *** denotes significance at the 1% level. Firm and year fixed effects and intercept terms are included in all specifications. The last column repeats the regression in the first column but excludes all patents that have not been renewed at least once.

Two factors that have been receiving some attention in the literature recently are licensing and litigation. They could influence a firm's incentives to do R&D, the value of a firm's patent stock, or a firm's decision to buy, keep, or sell patents. Licensing is an alternative vehicle for technology transfer. Additionally, litigation might affect a firm's decision to acquire, retain, or sell patents. Therefore, controls are introduced for litigation and licensing: columns 2–4 introduce the fraction of a firm's portfolio that is ever litigated, sector-level litigation intensity (defined as the fraction of litigated patents over total patents in that sector), and sector-level licensing intensity, respectively. Column 5 introduces all these controls at once. All of these alternative specifications show that the benchmark estimates in column 1 are remarkably robust. Last, some patents have little value. To control for this, the last column only includes those patents that were renewed at least once.¹¹ (Patents must be renewed, at a small fee, in their 3rd, 7th, and 11th years.) As can be seen, a lot of patents are not renewed and purging these patents increases somewhat the impact of the patent stock and distance-adjusted patent stock on the firm's market value. The story remains more or less the same, though, with the relative value of the first two regression coefficients more or less staying fixed. The gist of this section is summarized as follows:

FACT 2: A patent contributes more to a firm's stock market value if it is closer to the firm in terms of technological distance.

3.3.2. Patent Sale Decision and Patent–Firm Distance

Does the technological distance of a patent to the firm influence the decision to keep or sell it? In order to conduct this analysis, the indicator variable for whether a patent is sold or kept (= 1 if a patent is sold, = 0 if not) is regressed on a number of potentially related regressors, including the patent's distance to the initial owner. Table III reports the OLS regression results.

Using the full sample, column 1 of Table III indicates that a patent is more likely to be sold if it is more distant to the firm. The regression includes controls for the size of the patent portfolio of the firm, patent quality, year and firm fixed effects. The coefficient on the distance variable is statistically significant and positive. Considering the average number of patents sold ($\approx 16\%$) in the time period, the coefficient suggests that a perfectly mismatched patent is 13.1% ($\approx 0.0197/0.16$) more likely to be sold to another firm, rather than being kept. Recall also that the definition employed for a sale is quite conservative, in the sense that patent transfers due to mergers and acquisitions are not considered sales, even though the primary motive for these events might be the acquisition of patents. The results are in line with the intuition that a

¹¹Information on patent renewals is obtained from the USPTO's U.S. Patent Grant Maintenance Fee Events.

TABLE III
PATENT SALE DECISION (FULL SAMPLE WITH LITIGATION INTENSITY)^a

	Dependent Variable (= 1 if Sold, = 0 Otherwise)				
	(1)	(2)	(3)	(4)	(5)
Distance	1.972*** (0.078)	1.982*** (0.078)	1.980*** (0.078)	1.991*** (0.078)	2.398*** (0.091)
Tech-class litigation intensity	no	yes	no	yes	no
Patent litigation dummy	no	no	yes	yes	no
Renewed patents only	no	no	no	no	yes
Observations	2,547,881	2,547,881	2,547,881	2,547,881	1,976,964
R ²	0.42	0.42	0.42	0.42	0.44

^aThe dependent variable is multiplied by 100 for clarity. All specifications control for the number of citations that each patent received (using the HJT correction term), for the size of the firm's patent portfolio, and also for firm fixed effects. Year and intercept terms are included. The coefficients on the non-distance terms are not reported to save space and are available upon request. The last column repeats the regression in the first column of the panel but excludes all patents that have not been renewed once. Standard errors are in parentheses and *** denotes significance at the 1% level.

firm is more likely to sell patents that are not a good fit, rather than keeping them, due to the potential gains from trading the patent to a firm that might be better suited to exploit the embedded ideas commercially.

Column 2 controls for the litigation intensity of the technology class, while column 3 controls for the lifetime litigation status of the patent, and column 4 for both simultaneously. The association of distance to a patent sale is unaffected by the presence of these additional controls. Column 5 redoes the first regression but purges those patents that are not renewed at least once. The effect of distance is only slightly more pronounced. The results are also robust to inclusion of the licensing intensity of a sector—see Section B.6 in Empirical Appendix B for the results. The stark result is that adding additional controls or restricting the sample do not weaken the link between distance and patent sale; to the contrary, they make it more pronounced. The implications of litigation and licensing on the market for ideas will be explored in more detail in Section 6.

FACT 3: *A patent is more likely to be sold the more distant it is to a firm.*

The primary motivation behind considering patent distance as a likely determinant of patent sale decisions is the potential gains from trade that arise if the patent can be sold to a firm that can use it better, which in expectation yields more profits. If this intuition is correct, the distance between the owner firm and the patent is expected to decrease after a patent is sold. Let $d(p, f_b)$ denote the distance of the patent to the buyer firm, and $d(p, f_s)$ to the seller firm. Next, the change in distance, $d(p, f_b) - d(p, f_s)$, is computed. This difference is -0.152 in 1980, the beginning of the sample, with a standard error of 0.049 .

What this shows is that conditional on a patent sale, the distance between a patent and its owner is significantly decreased. In other words, the mismatch between the idea and the firm owning it is reduced. The effect is economically large. Considering that the average measure for distance is 0.481, the average reduction in distance is approximately 32% ($\simeq 0.152/0.481$) of the average distance. The average distance reduction in the whole sample is 16% and this number goes up to as high as 49% in 2006, which is the end year of the sample.

FACT 4: *A patent is technologically closer to the buying firm than to the selling firm.*

3.4. Tacking on a Market for n -Type Patents

To append a market for n -type ideas onto the model, recall that a firm obtains an n -type idea with probability p . This can arise in one of two ways: either the firm develops an n -type idea or it purchases one. Let a firm that develops an n -type idea sell it with probability p_s . Likewise, assume that a firm that fails to come up with an n -type idea will purchase one with probability p_b . Suppose that the market for n -type ideas clears instantaneously every period. This implies that $pp_s = (1 - p)p_b$, so that $p_b = p_s p / (1 - p)$. Adding a market for n -type patents onto the above structure does *not* alter the model's solution for a symmetric balanced growth path. This is discussed further in Theory Appendix A, Section A.2.

In the U.S. data, the distance between a patent and its owner's line of business shrinks on average upon a sale; that is, a patent is closer to the buyer than the seller. This is not true empirically for all patent sales. The presence of n -type patents helps the model better capture Fact 4. It is easy to deduce that, on average, the distance between a d -type patent and its owner would contract in the model by $[1/(1 - x_a)] \int_{x_a}^1 x dx - [1/X(x_k)] \int_0^{x_k} x dX(x)$, since a non-innovating business buys if $x > x_a$ and an innovating firm sells when $x < x_k$. The average distance between an n -type patent and its owner would contract in the model by $\int_0^1 x dx - \int_0^1 x dX(x)$.¹² This is smaller than the number for d -type patents, because $[1/(1 - x_a)] \int_{x_a}^1 x dx > \int_0^1 x dx$ and $[1/X(x_k)] \int_0^{x_k} x dX(x) < \int_0^1 x dX(x)$. Thus, the presence of a market for n -type patents operates to reduce the average shrinkage in distance upon sale between a patent and its owner.

4. CALIBRATION

In order to simulate the model, values must be assigned to the various parameters. There are sixteen parameters to pick: β , ε , κ , λ , δ , σ , γ_d , χ , ρ , μ ,

¹²The distance between an n -type patent and its owner has no real effect; that is, the technology class for an n -type patent is just a label.

η , ω , γ_n , p , p_s , and $\text{STD}(e')$. A distribution for $X(x)$ needs to be provided as well. As is standard in macroeconomics, some of the parameter values are chosen on the basis of a priori information, while others are determined internally using a minimum distance estimation routine. By selecting some parameters using a priori information, the size of the calibration/estimation procedure is reduced. This is important because undertaking calibration/estimation is problematic when there is a large number of parameter values. For the most part, there is either a consensus about what the appropriate values for these parameters are, or the U.S. data speak directly to them. The selection of parameter values on the basis of a priori information is now discussed.

4.1. *The Use of a priori Information*

1. *Capital's and labor's shares of income*, κ and λ . In line with Corrado, Hulten, and Sichel (2009) estimates from the U.S. National Income and Product Accounts, capital's and labor's shares of incomes, κ and λ , are set to 25 and 60%. This implies that the profit parameter, as represented by ζ , accounts for the remaining 15%. This is a fairly typical value used in the macroeconomics literature, as is discussed in Guner, Ventura, and Xu (2008).

2. *Depreciation rate for capital*, δ . The depreciation rate of capital is chosen to be 6.9%. This is consistent with the U.S. National Income and Product Accounts.

3. *Survival rate for a patent*, σ . In the United States, a patent lasts for 17 years. Hence, $\sigma = 1 - 1/(1 + 17)$.

4. *CRRRA parameter*, ε . This parameter is taken to be 2, the midpoint between the various estimates reported in Kaplow (2005). This is a common value used in macroeconomics.

5. *Long-run interest rate*. A reasonable value for the long-run interest rate in the United States is 6%—see Cooley and Prescott (1995). Now, the long-run growth rate for the United States is 2%. Given the value for the economy's long-run growth rate, $\mathbf{g}^{\zeta/(\zeta+\lambda)} = 1.02$, and the coefficient of relative risk aversion, $\varepsilon = 2$, the discount factor, β , is then uniquely pinned down using the equation $\beta = r \mathbf{g}^{\varepsilon \zeta/(\zeta+\lambda)}$ —see (22). This is standard procedure for a growth model.

6. *The empirical distribution for the proximity of patent to a firm's technology class*. The empirical distance distribution for the United States displayed in Figure 3, for $\iota = 2/3$, is used for the analysis. Define a measure of propinquity (or closeness) between a patent p and a firm f by $c_i(p, f) \equiv 1 - d_i(p, f)$, where $d_i(p, f)$ is given by (24). The density associated with $c_i(p, f)$ is used for $X(x)$. This amounts to just a simple change in units on the horizontal axis in Figure 3. Assume that x is distributed uniformly within each of the ten bins of the histogram. (There is an additional mass point at one.) One might think that a firm will try to invent ideas that are close to its line of business. The calibration strategy forces the propinquity of ideas to the inventor's line of business in the model to be congruent with the U.S. data.

7. *R&D cost elasticity*, ρ . In order to estimate the elasticity of the R&D cost function, the cost function in the model is inverted to obtain a production function. Then, a regression is run using Compustat data to determine the parameter value, ρ , where the output of the R&D production function is proxied for by citation-weighted patents.

8. *Bargaining power*, ω . The bargaining powers of buyers and sellers are chosen to be equal. This assumption is often imposed in macroeconomic models using Nash bargaining. Unfortunately, there does not seem to be a good way to identify a value for this parameter, either using a priori information or through the calibration/estimation procedure discussed below. Due to the presence of a spillover externality in (2), the Hosios condition will not necessarily lead to an efficient matching equilibrium.

Therefore, values for the parameters β , ε , κ , λ , δ , σ , ρ , and ω are imposed using a priori information in line with points (1), (2), (3), (4), (5), (7), and (8) without having to solve the model. The distribution $X(x)$ is constructed in line with point (6).

4.2. Minimum Distance Estimation

Values for the remaining parameters, χ , μ , γ_d , η , γ_n , p , p_s , and $\text{STD}(\ln e')$, must be assigned. This is done by minimizing the sum of the squares between some data targets, discussed below, and the model's predictions for these targets. The model is highly nonlinear in nature. Computing the solution to the model essentially involves solving a system of nonlinear equations, as is discussed in Theory Appendix A, Section A.1. Therefore, it is not the case that a particular parameter is identified uniquely by a particular data target. By computing the Jacobian of the system, the influence of each parameter on the data targets can be gauged. The presentation below uses this Jacobian and other features of the framework to discuss, in a heuristic fashion, how the parameters are identified. The Jacobian is presented in Section B.7 of Empirical Appendix B. The data targets are listed in (1) to (7) below. Targets (1) to (5) are discussed now.

1. *Long-run growth in output*. In the United States, output grew at about 2% per year over the postwar period. Intuitively, one would expect the parameter γ_d , which governs how d -type innovations enter the law of motion for a firm's productivity growth (2), should play an important role in determining this. The same is true for the n -type patent parameters, γ_n and p . The Jacobian confirms that these parameters have a positive impact on growth—see Section B.7 for more detail. The term for the d -type patents, or γ_d , dominates the others. The parameter governing the cost of R&D, χ , has a negative and smaller effect on growth.

2. *The ratio of R&D expenditure to GDP*. U.S. expenditure on research and development is about 2.91% of GDP—see [Akcigit, Hanley, and Serrano-Velarde \(2014\)](#). What parameters influence this ratio? Again, the parameter γ_d

governing the productivity of d -type patents is very important. It increases this ratio because the payoff from R&D rises with γ_d . Not surprisingly, the R&D cost parameter, χ , has a bearing here, because it directly governs the cost of innovation, as can be seen from (4). Last, the n -type patent parameters, γ_n and p , are negatively associated with this ratio. They increase GDP growth without the need to do R&D.

3. *Fraction of patents sold.* About 16% of patents are sold in the United States, as catalogued in Table I. The parameters governing the matching function, μ and η , control how easy it is to sell a d -type patent. They are important in determining this ratio. The parameters, p and p_s , regulating the arrival and sales rates for n -type patents are also important, although the dependence here is of a mechanical nature.

4. *Duration until a sale.* The entire empirical frequency distribution for the duration of a sale is targeted—see Figure 4.¹³ In particular, the calibration procedure tries to minimize the sum of the squared differences between the empirical distribution and its analogue for the model. It takes about 5.34 years on average to sell a patent. The coefficient of variation around this mean is 0.84. So, there is considerable variation in sale duration. The parameters governing the matching function, μ and especially η , are obviously central here. This




FIGURE 4.—Sales duration distribution, data and model.

¹³Different criteria can be used for dating when an idea is born. One could use the application date instead of the grant date since some patents are sold before they are granted. An alternative would be to use the first time that another inventor builds on this invention (as measured by the first citation that a patent receives). This reflects the time that it took for others to learn about the idea. Last, it is possible that excluding more recent observations might prevent the confounding effects of a potential truncation bias. This occurs because patents toward the end of sample have less time to be sold. Repeating the analysis using these three new sale duration distributions does not change the main findings.

can be seen from equation (27) in Section A.1 of Theory Appendix A, which specifies the odds that a patent agent will find a buyer. These parameters also influence the spread in duration.

5. *Distance reduction upon sale—all patents.* Section 3.3.2 presents an estimate (-0.152) from the micro data on the average difference between a buyer's and seller's technological propinquity for a patent.¹⁴ This estimate is targeted and helps to discipline the relative importance of d - and n -type patents. As is discussed in Section 3.4, the arrival rate of n -type ideas and the probability of selling them, or p and p_s , are central here. They operate to reduce the observed amount of distance reduction since the sale of these patents does not depend upon technological propinquity. This is shown by the Jacobian of the system. Additionally, the parameters of the matching function, μ and η , influence the model's ability to hit this target. More efficient matching implies a larger reduction in distance.

4.2.1. Indirect Inference

The data targets (6) and (7) discussed below derive from the firm-level panel-data regressions presented in Section 3. As was mentioned, computing the equilibrium solution for the model essentially involves solving a system of nonlinear equations, as Theory Appendix A, Section A.1, makes clear. Undertaking the indirect inference involves an additional step. Here a Monte Carlo simulation is undertaken on a panel of 30,000 firms for 30 periods (to replicate the number of periods in the data). This is used to estimate the panel-data regressions analogues for the model that correspond with the ones estimated from the U.S. data, which are presented in Table II.

6. *Relative strength of the patent stock versus the distance-adjusted patent stock on a firm's market value.* This is estimated from the micro data—Table II, column 1. It is measured by the ratio of the coefficient on log distance-adjusted patent stock to log patent stock. This target plays a significant role in identifying the size of the distance related term, γ_d , relative to the non-distance related ones, γ_n and p , in the law of motion for productivity (2). The former has a positive impact on this ratio, while the latter have negative ones. The matching function parameters, μ and η , also have an influence on this target because they affect the value of a d -type patent. Similarly, so does the cost of doing R&D, χ . Last, the probability of selling an n -type patent, conditional upon its arrival, p_s , affects this statistic. The higher the likelihood that an n -type patent is sold, and therefore that it is not used in production, the less impact it will have on a firm's market value. This results in d -type patents mattering more for market value relative to n -type ones; again, the detail is in the Jacobian presented in Section B.7.

¹⁴The quantitative results do not change in a material way when the mean of the averages over all years in the sample is used instead.

7. *Relative strength of the patent stock versus employment on a firm's market value.* This, too, is estimated from the micro data—Table II, column 1. It is measured as the ratio of the coefficient on log patent stock to the coefficient on log employment. The (inverse of this) ratio can be thought of as measuring the impact of an increase in the patent stock on employment, holding fixed the firm's market value. In the model, there are two reasons a firm's market value may rise relative to other firms. Its long-run productivity, z' , may have increased relative to average long-run productivity, \mathbf{z}' , or it may have realized a favorable value for the temporary production shock, e' . This ratio identifies the standard deviation of the firm-specific idiosyncratic production shock, $\text{STD}(\ln e')$. Without the e' shock, employment would be a perfect predictor of relative productivity, z'/\mathbf{z}' .¹⁵ Introducing the e' shock breaks this one-to-one correspondence. The parameter $\text{STD}(\ln e')$ has no impact on the other data targets. The parameter γ_n governing the productivity of n -type patents also affects this ratio. As γ_n rises, employment becomes a better predictor of a firm's market value, so it impinges on this ratio in a negative way. An increase in γ_d does not work the same way as it results in more d -type ideas, which makes the patent stock a better predictor of market value.

To highlight a central aspect of the calibration procedure, note that a key goal of this research here is to quantify the importance of the patent market for eliminating the misallocation of ideas across producers. Two considerations come into play: the importance of technological propinquity between a patent and a producer (or γ_d) and the efficiency of the market for ideas (or η). A low volume of patent sales could occur either because technological propinquity is not very important (but the patent market is still efficient) or because the patent market is inefficient (but technological propinquity is important). The above micro data are used to identify both of these channels. At the risk of sounding repetitive, the firm market-value regressions in Table II are used to speak to the size of γ_d . Since firm fixed effects are included in these regressions, there is a strong sense in which changes in the distance-adjusted patent stock are being tied to firm market value. Therefore, reproducing similar regression results using the model-generated data (in particular, the relative size of the coefficients on the log distance-adjusted patent stock to the log patent stock) helps identify γ_d . Matching up the model's output with the micro data on the fraction of patents sold, average sale duration, and the difference between the buyer's and seller's technological propinquinities pins down η . The efficiency of the market for ideas plays a very important part in the analysis and is analyzed in detail in Section 5.

It is well known that patents show big differences in terms of their qualities, which could also affect their sales. A reasonable belief might be that a small fraction of patents are highly valuable while the median one is not. To take

¹⁵To see this, substitute (20) into (19), while making use of the definition in (18), to get $l = e' z' / \mathbf{z}'$.

TABLE IV
PARAMETER VALUES USED IN THE BASELINE SIMULATION

Parameter Value	Description	Identification
$\beta = 0.98$	Discount factor	A priori information
$\varepsilon = 2.00$	CRRA parameter	A priori information
$\kappa = 0.25$	Capital's share	A priori information
$\lambda = 0.60$	Labor's share	A priori information
$\delta = 0.07$	Depreciation rate	A priori information
$\sigma = 0.94$	Patent survival rate	A priori information
$\gamma_d = 0.25$	Distance-related productivity	Calibration/estimation
$\chi = 0.83$	Cost of R&D	Calibration/estimation
$\rho = 3.00$	R&D cost elasticity	A priori information
$\mu = 0.50$	Matching function, exp	Calibration/estimation
$\eta = 0.09$	Matching function, const	Calibration/estimation
$\omega = 0.50$	Bargaining power	Imposed
$X(x)$	Proximity distribution	A priori information
$\gamma_n = 0.18$	Non-distance related productivity	Calibration/estimation
$p = 0.17$	Pr(n -type idea)	Calibration/estimation
$p_s = 0.47$	Pr(sell n -type patent arrival)	Calibration/estimation
$\text{STD}(\ln e') = 0.07$	Production shock, std	Calibration/estimation

quality heterogeneity into account, all regressions control for patent citations as a proxy for patent quality. So, the empirical analysis attempts to purge concerns about patent quality from the stylized facts.¹⁶

The upshot of the calibration procedure is displayed in Tables IV and V. Figure 4 shows, for both the data and model, the frequency distribution over the duration for a sale. As can be seen, it appears to be harder to affect a sale in data than in the model.

5. FINDINGS

The importance of a market for patents will be gauged now. There are two sources of inefficiencies in the model. The first one is the usual knowledge externality. Each single innovation raises the aggregate knowledge stock in society, which benefits the future generations that stand on the shoulders of former giants through z in (2). The second source of inefficiency emerges due to *matching frictions*, which is of particular interest here. To analyze the latter, various experiments that change the efficiency of the market for d -type

¹⁶Alternatively, one could introduce quality into the model. In particular, every idea could have a quality component drawn from some distribution. Now, the decisions to buy and sell patents would be a function of distance and quality (in addition to the aggregate state variable). Perhaps the distribution governing quality could be mapped into the empirical distribution for patent citations. Doing this would significantly complicate the analysis, but could be a fruitful avenue for future research.

TABLE V
CALIBRATION TARGETS^a

Target	U.S. Data	Model
Long-run growth in output	2.00%	2.08%
Ratio of R&D expenditure to GDP	2.91%	1.96%
Fraction of patents that are sold	15.6%	16.6%
Average duration until a sale (fit entire distribution)	5.48 yrs	6.28 yrs
Sale duration, c.v. (fit entire distribution)	0.84	0.71
COEF(dist-adj pat stock)/COEF(pat stock)	-0.511	-0.590
COEF(pat stock)/COEF(empl)	0.054	0.054
$d(p, f_s) - d(p, f_b)$, all sold	0.152	0.165

^aIn the calibration, the full sales duration distribution (17 points) is targeted. The above table just reports the mean and the coefficient of variation for this distribution as summary measures.

patents will be entertained. The efficiency of the market for d -type patents is increased in stages. First, the market is shut down by setting the meeting rate to zero. Then, an experiment is performed where the meeting rate for matches is allowed to rise. While it may be easier for buyers and sellers to meet now, a seller's idea may still not be well suited for the buyer. The next experiment considers a situation where patent agents can find buyers who are perfect matches for the ideas that they are selling. So, there is no mismatch between buyers and sellers on the patent market. Still, innovating firms produce d -type ideas that are not ideally suited for their own businesses and this injects a friction into the analysis. A patent that is not incorporated into an innovator's production process will only have a finite life on the market. Additionally, it may take time to find a buyer. The final experiment focuses on the case where innovating firms produce ideas that are tailored toward their own production activity. Here ideas are perfectly matched with the developer. The change in welfare from moving from one environment to another is calculated. The metric for comparing welfare is discussed now.

5.1. Welfare Comparisons

Consider two economies, namely A and B , moving along their balanced growth paths. Aggregate consumption, the gross growth rate, and aggregate productivity for economy A are represented by \mathbf{c}^A , \mathbf{g}^A , and \mathbf{z}^A . Similar notation is used for country B . To render things comparable, start each country off from the same initial position where $\mathbf{z}^A = \mathbf{z}^B = 1$. Now, the levels of welfare for economies A and B are given by

$$W^A = \sum_{t=1}^{\infty} \beta^{t-1} \frac{(\mathbf{c}_t^A)^{1-\varepsilon}}{1-\varepsilon} = \frac{(\mathbf{c}_1^A)^{1-\varepsilon}}{(1-\varepsilon)[1-\beta(\mathbf{g}^A)^{1-\varepsilon}]},$$

and

$$W^B = \frac{(\mathbf{c}_1^B)^{1-\varepsilon}}{(1-\varepsilon)[1-\beta(\mathbf{g}^B)^{1-\varepsilon}]},$$

where \mathbf{c}_1^A and \mathbf{c}_1^B are the time-1 levels of consumption in economies A and B . How much would initial consumption in economy A have to be raised or lowered to make people have the same welfare level as in economy B ? Denote the fractional amount in gross terms by α (which may be less than 1). Then, α must solve

$$\frac{(\alpha \mathbf{c}_1^A)^{1-\varepsilon}}{(1-\varepsilon)[1-\beta(\mathbf{g}^A)^{1-\varepsilon}]} = W^B,$$

so that

$$\alpha = (W^B / W^A)^{1/(1-\varepsilon)}.$$

This welfare measure is used in all experiments.

5.2. Varying the Contact Rate for Matches, η

The patent market mitigates the initial misallocation of ideas. Still, it takes time to sell a patent as the patent agent may not be able to find a buyer. To understand how this friction in matching affects the economy, it is useful to examine the relationship between the scale factor for the matching function, η , and several key variables.¹⁷ Figures 5 and 6 summarize the results.




FIGURE 5.—The impact of an increase in the contact rate on duration, innovation, growth, and welfare.

¹⁷The relationship between \mathbf{g} and η is highlighted in the simplified model developed in Akcigit, Celik, and Greenwood (2015, Appendix 12).




FIGURE 6.—The impact of an increase in the contact rate on the price for the innovator, slackness, and the cutoffs.

The market for d -type patents is shut down when $\eta = 0$. When there is no market, the equilibrium growth rate goes down to 2.02% from its benchmark value of 2.08%. Shutting down the market results in a welfare reduction of 1.18% in consumption equivalent terms, which is quite sizable. As the contact rate, η , rises, it becomes easier to find a buyer for a patent, *ceteris paribus*. This is reflected in a drop in the length of time that it takes to find a buyer, as the right panel of Figure 5 illustrates. The price that an innovating firm receives for a patent, q , rises accordingly—see the left panel of Figure 6. As the price moves up, an innovating firm becomes choosier about the patents that it will keep. Figure 6, right panel, illustrates how an innovator's cutoff for selling, x_k , rises with η . (Recall that better patents are associated with higher values for the propinquity metric.) Similarly, buyers become pickier about the patents that they will purchase so that x_a moves up with η .

The rate of innovation, \mathbf{i} , does not change much. It falls as η starts to rise since the consequences of failing to innovate are now lessened, because it will be easier for a firm to buy a patent. The high price for patents begins to spur innovation at higher levels of η . Market slackness, n_a/n_b , has an interesting \cap shape, which is displayed in the left panel of Figure 6. When $\eta = 0$, the patent market is essentially closed, as no innovators will want to sell their ideas. The number of prospective buyers is $1 - \mathbf{i}$. As η starts to rise, so does the number of innovators that want to sell their ideas. This increases the flow of new patents into the patent market and results in n_a/n_b moving upwards. As the rate of innovation, \mathbf{i} , declines, the number of prospective buyers, $1 - \mathbf{i}$, rises. This force operates to reduce n_a/n_b . Additionally, as the contact rate increases, the market becomes more efficient. It is easier for a seller to find a buyer, *ceteris paribus*. This works to reduce the stock of sellers.

Growth increases along with efficiency in matching, despite the reduction in the number of new ideas—see the left panel of Figure 5. So does welfare. If the efficiency of the market was at its extreme (the minimum value for η

that results in all buyers meeting a patent agent with probability 1), growth would go up to 2.46% and welfare would be 5.97% higher than the calibrated economy. The upshot is that the market for patents plays an important role in the economy.

5.3. Perfectly Directed Search

A second source of inefficiency in the model is the random search technology used in the d -type patent market. In the baseline model, conditional upon a meeting between a buyer and a patent agent, the propinquity of the idea to the firm is drawn from a uniform distribution. Instead imagine a perfectly directed search structure, where patent agents are able to target the segment of the economy that exactly matches the patent they want to sell. In such a case, whether or not a patent agent meets a buyer is still a probabilistic event governed by the matching function. The propinquity between the patent and the buying firm would be nonstochastic and equal to unity; in other words, a perfect match. The level of welfare in this alternative economy is 1.94% higher than in the baseline one. The output growth rate increases slightly from 2.08 to 2.19%, despite a small decline in the innovation rate. The fraction of all patents sold increases from 16.6 to 19.9%. Last, a decomposition of growth reveals that the fraction of growth due to all patents sold moves up from 18.9 to 26.6%.¹⁸ Table VI summarizes the results (where the baseline model is labeled BM and PDS refers to the perfectly directed search structure).

TABLE VI
THOUGHT EXPERIMENTS^a

	BM	PDS	PDSwHC	PI
Output growth rate, %, $(\mathbf{g}^{\xi/(\xi+\lambda)} - 1) \times 100$	2.08	2.19	3.05	3.38
Innovation rate, \mathbf{i}	0.58	0.56	0.57	0.61
Welfare gain, $\alpha - 1$	0.00	0.02	0.14	0.18
Fraction of all patents sold	0.17	0.20	0.68	0
Growth from all patents sold	0.19	0.27	0.73	0

^aThe first column of results is for the baseline model (BM). Perfectly directed search (PDS) is shown in the second column where a patent sold is a perfect match for the buyer ($x = 1$). In the third column (PDSwHC), there is perfectly directed search, plus there is a high contact rate between patent agents and buyers. All innovating firms draw the perfect idea ($x = 1$) in the last column (PI). The figures in the first row (only) are in percent.

¹⁸Section A.1 shows that $\mathbf{g} - \mathbf{1} = \gamma_d \mathbf{i} \int_{x_k}^1 x dX(x) + \gamma_d(1 - \mathbf{i})m_b(\frac{n_a}{n_b}) \int_{x_a}^1 x dx + \gamma_n \mathbf{p}$, where \mathbf{i} is the aggregate rate of innovation. Note that there are three terms on the right side. The first term can be used to measure the contribution to growth from the distance related ideas that firms keep, the second from the ones that they sell. The third term gives the growth arising from non-distance related ideas. This term can be further decomposed as $\gamma_n \mathbf{p} = \gamma_n[\mathbf{p}(1 - \mathbf{p}_s) + (1 - \mathbf{p})\mathbf{p}_b]$, where the first term in brackets gives the contribution from non-distance related patents kept and the second from the ones sold.

5.3.1. Perfectly Directed Search With a High Contact Rate

Now, redo the above experiment with perfect directly search while also using a high contact rate for matches.¹⁹ The results are reported in Table VI (under the column labeled PDSwHC). Output growth is now much higher at 3.05%, even though innovation is slightly lower than in the baseline model. This reflects a reduction in misallocation. As can be seen, now most patents are sold. Economic welfare is 14.3% higher.²⁰

Figure 7 gives the upshot from the experiments that have been conducted so far. It shows how the cumulative distribution function for the propinquity of new ideas to firms, or for x , changes across the various experiments. First, firms in the U.S. data produce ideas that are not well suited for their own lines of business, as can be seen from the distribution labeled “Empirical.” (Recall that a higher value for $x \in [0, 1]$ indicates that an idea is better suited for the firm’s business activity.) In the baseline model, a firm is free to sell such an idea. A firm that fails to innovate can try to buy one from another firm. This leads to a better distribution of ideas, as is reflected in the distribution function for the baseline model after transactions on the market for patents have




FIGURE 7.—Misallocation of ideas. The graph plots the cumulative distribution functions for x . A higher value for x , measuring propinquity, implies that an idea is better suited for a firm.

¹⁹The contact rate, η , is set high enough that all buyers meet a patent agent with probability 1.

²⁰This large welfare gain derives solely from the large increase in growth, $g^{\varepsilon/(\varepsilon+\lambda)}$, given the assumed form of preferences over consumption, as can be gleaned from Section 5.1. That is, if there is a large increase in growth, then this form of preferences will always show a large increase in welfare (when $\varepsilon = 2$, which is a standard value).

been consummated. The distribution function for the baseline model stochastically dominates, in the first-order sense, the empirical distribution. When the contact rate for matching is high, it is relatively easy to consummate a patent sale. The distribution for x improves—see the histogram labeled “High Contact Rate,” which stochastically dominates the one for the baseline model. Of course, if search could be perfectly directed, things would be better still—“High Contact w Directed Search,” which stochastically dominates all other distributions.

Note that not all firms sell their patents, even though they are not perfectly matched with their ideas. This occurs because there are still some frictions left in the patent market. First, there are more sellers than buyers on the market, so not all patents will be immediately sold. Second, patents have a finite life on the market and hence suffer some depreciation. Both these factors imply that the price at which a firm can sell a patent, q , will be less than what it is worth to a perfectly matched firm.

5.4. *Removing the Misallocation of Ideas*

The central inefficiency in the framework derives from the fact that firms develop ideas that are imperfect matches for their own production processes. The presence of a market for patents mitigates this problem. Suppose that an innovating firm comes up with a d -type idea that is always a perfect match for its production process. That is, let each innovating firm always draw $x = 1$. This case is summarized in Table VI (under the column labeled PI). In this situation, the economy could increase its growth rate from 2.08 to 3.38%, a big jump. Welfare would increase by 17.8%. This illustrates that the frictions arising from mismatches in innovation are large.

6. QUANTITATIVE EXTENSIONS: LICENSING AND LITIGATION

When it comes to technology transfer and the market for ideas, two important concerns about patenting and the market for ideas deserve additional attention, namely licensing and litigation. Licensing provides an additional mechanism for transferring ideas. By limiting attention to patent sales, a fear might be that the analysis overstates the amount of misallocation in the market for ideas. A firm may buy or keep a patent to prevent litigation. This does not increase the firm’s productivity in a technological sense. Hence, the value of patents for a firm’s productivity may be overestimated.

6.1. *Licensing*

Arora and Ceccagnoli (2006) reported that licensing intensity in the United States is around 5%. The goal here is to understand the quantitative implications of licensing in the current setting. Zuniga and Guellec (2009) conducted

TABLE VII
RESULTS WITH LICENSING AND LOW-LITIGATION SECTORS^a

Panel A. Calibration Targets				
	Licensing		Low-Litigation	
	U.S. Data	Model	U.S. Data	Model
Growth in output, %	2.00	2.03	2.00	2.10
Ratio of R&D expenditure to GDP, %	2.91	1.81	2.91	1.98
Fraction of ideas that are sold, %	20.6	20.5	16.4	17.6
Average duration until a sale, yrs	5.48	6.05	5.94	6.35
Sale duration, c.v.	0.84	0.72	0.78	0.70
COEF(dist-adj pat stock)/COEF(pat stock)	-0.511	-0.607	-0.568	-0.596
COEF(pat stock)/COEF(empl)	0.054	0.057	0.052	0.050
$d(p, f_s) - d(p, f_b)$, all sold	0.152	0.161	0.136	0.143
Panel B. Impact of Shutting Down the Market for Ideas ($\eta = 0$)				
	Benchmark	Licensing	Low-Litigation	
Δ in growth (percentage pt)	-0.06	-0.07	-0.06	
Δ in welfare, %	-1.18%	-1.40%	-1.12%	

^aResults for both the data and model when ideas can be also transferred via licensing and when the analysis is restricted to low-litigation sectors.

a survey on firms that license out their patents and analyzed the obstacles to licensing. The most frequent problem reported by firms was that “identifying (a) partner is difficult.” This shows that search frictions, which are highlighted in the model of the patent market developed here, seem to apply to the licensing market as well. Licensing could have many other purposes than pure technology transfer, such as deterring entry. To the extent that licensing is used as a substitute arrangement for a patent sale, the previous analysis might have underestimated the liquidity in the market for ideas and generated too much search frictions. In order to take this substitutability into account, assume that all the licensing arrangements are for the purpose of technology transfer. Hence in what follows, assume that the overall turnover in the market for ideas is $20.6\% = 15.6\% + 5\%$. The model is recalibrated and simulated using this number. Table VII reports the results. The model matches the data well when it is recalibrated to allow for a larger number of ideas to be transferred. Not surprisingly, a shutdown in the market for ideas leads now to a bigger welfare loss (1.40 versus 1.18%). As before, the reduction in growth is still small, but slightly higher (a loss of 0.07 versus 0.06 percentage points). Again, the small loss in growth is due to the fact that the rate of innovation rises when the market for ideas is closed, as was shown earlier in the right panel of Figure 5.

6.2. *Litigation*

Patent litigation could also lead to patent sales for reasons not necessarily related to technology transfer (Galasso, Schankerman, and Serrano (2013)). To begin with, it is useful to get a sense of the share of patents that are ever litigated in the sample employed here. Using the Derwent and Lex Machina databases, Table I shows that about 1.0% of patents involve litigation. Furthermore, when patents that are both ever litigated and ever sold during their lifetime are considered, the share drops down to 0.3%. Hence, among sold patents, only 2% ($= 0.3/15.6$) are ever litigated. Given these small shares, it may seem unlikely that litigated patents could have a major impact on the quantitative results.

As Galasso, Schankerman, and Serrano (2013) emphasized, however, the threat of litigation might be very important in the sale decision, even if in practice, few litigations are actually observed. In order to exclude this potential channel, the analysis is redone, focusing exclusively on sectors with very low litigation intensity. All the micro data targets are recalculated using patent and firm observations that have a litigation intensity below the mean of the pertinent sample; the targets for the U.S. growth rate and R&D expenditures to GDP remain the same. Indeed, sectors have a lot of heterogeneity in terms of the litigations observed, and a sector's litigation intensity might be a good indicator for the propensity of a given patent to be litigated.

Table VII presents the new estimates and the welfare gain from the market. Note that the affected data targets change only slightly. These changes occur from restricting the micro data to the low-litigation sectors. The model still fits very well. The welfare loss for shutting down the market for ideas is now a bit smaller (1.12 versus 1.18%). The upshot is that focusing on low-litigation sectors does not affect the analysis in a material way.

7. CONCLUSIONS

A model of the market for patents is developed here. Each period, a firm conducts research and development. This R&D process may spawn new ideas. Some of the ideas are useful for a firm's line of business; others are not. A firm can patent and sell the ideas that are not. The fact it can sell ideas provides an incentive to engage in R&D. Likewise, firms that fail to innovate can attempt to buy ideas. This allows a firm to grow its business. This reduces the incentive to do R&D. The efficiency of the patent market for matching ideas with firms has implications for growth. These are examined here.

The empirical analysis, drawing on the NBER-USPTO patent grant database and patent reassignment data available from Google Patents Beta, establishes five useful facts. First, somewhere between 15% and 20% of patents are sold. Second, it takes, on average, 5.48 years to sell a patent. Third, a firm's patent stock contributes more to its market value the closer it is to the firm in terms of average technological distance. Fourth, a patent is more likely to be sold the

more distant it is to a firm's line of business. Fifth, when a patent is sold, it is closer to the buyer's line of business than to the seller's. The empirical analysis attempts to control for licensing and litigation. These five facts suggest that a market for patents may play an important role in correcting the misallocation of ideas across firms. It may also influence a firm's R&D decision.

The developed model is calibrated to match several stylized facts characterizing the U.S. data, such as the postwar rate of growth, the ratio of R&D spending to GDP, the fraction of patents sold, the empirical sale duration distribution, and the reduction in distance between a patent and its owner upon a sale. Additionally, some micro-level facts from panel data regressions are targeted using an indirect inference strategy. Specifically, the importance of distance in a firm's patent portfolio for determining the firm's market value is zeroed in on. The value of a market for selling patents is then assessed. This is done by conducting a series of thought experiments where the market is first shut down and then the efficiency of the patent market is increased successively. The efficiency of this market is important for economic growth and welfare.

The new USPTO patent reassignment data open new and exciting directions for future research on innovation and technological progress. One direction is the analysis of optimal patent policy that not only considers the monopoly distortions and innovation incentives, but also takes into account the possibility of trading ideas through patents. Another direction is the analysis of firm dynamics when patents are not only produced in-house, but also purchased from others. Finally, the role of financial frictions is also a new and important channel that could impact the (mis)allocation of ideas. These are all very exciting and important aspects of technological progress that await further research.

REFERENCES

- ACEMOGLU, D., U. AKCIGIT, N. BLOOM, AND W. R. KERR (2013): "Innovation, Reallocation and Growth," Working Paper 18993, NBER. [948]
- AGHION, P., AND P. HOWITT (1992): "A Model of Growth Through Creative Destruction," *Econometrica*, 60, 323–351. [947]
- AKCIGIT, U., M. A. CELIK, AND J. GREENWOOD (2015): "Buy, Keep or Sell: Economic Growth and the Market for Ideas," RCER Working Paper 593, University of Rochester. [944,949,959,975]
- (2016): "Supplement to 'Buy, Keep, or Sell: Economic Growth and the Market for Ideas'," *Econometrica Supplemental Material*, 84, <http://dx.doi.org/10.3982/ECTA12144>. [944]
- AKCIGIT, U., D. HANLEY, AND N. SERRANO-VELARDE (2014): "Back to Basics: Basic Research Spillovers, Innovation Policy and Growth," Unpublished Manuscript, University of Pennsylvania. [969]
- ARORA, A., AND M. CECCAGNOLI (2006): "Patent Protection, Complementary Assets, and Firms' Incentives for Technology Licensing," *Management Science*, 52, 293–308. [979]
- BENHABIB, J., J. PERLA, AND C. TONETTI (2014): "Catch-up and Fall-Back Through Innovation and Imitation," *Journal of Economic Growth*, 19, 1–35. [947]
- CHIU, J., C. MEH, AND R. WRIGHT (2011): "Innovation and Growth With Financial, and Other, Frictions," Working Paper 17512, NBER. [947]

- COHEN, W. M., R. R. NELSON, AND J. P. WALSH (2000): "Protecting Their Intellectual Assets: Appropriability Conditions and Why U.S. Manufacturing Firms Patent (or Not)," Working Paper 7552, NBER. [960]
- COOLEY, T. F., AND E. C. PRESCOTT (1995): "Economic Growth and Business Cycles," in *Frontiers of Business Cycle Research*, ed. by T. F. Cooley. Princeton, NJ: Princeton University Press, 1–38. [968]
- CORRADO, C., C. HULTEN, AND D. SICHEL (2009): "Intangible Capital and U.S. Economic Growth," *Review of Income and Wealth*, 55, 661–685. [968]
- FIGUEROA, N., AND C. J. SERRANO (2013): "Patent Trading Flows of Small and Large Firms," Working Paper 18982, NBER. [948]
- GALASSO, A., M. SCHANKERMAN, AND C. J. SERRANO (2013): "Trading and Enforcing Patent Rights," *RAND Journal of Economics*, 44, 275–312. [960,981]
- GORT, M. (1962): *Diversification and Integration in American Industry*. Princeton, NJ: Princeton University Press. [948]
- GUNER, N., G. VENTURA, AND Y. XU (2008): "Macroeconomic Implications of Size-Dependent Policies," *Review of Economic Dynamics*, 11, 721–744. [948,968]
- HALL, B. H., A. B. JAFFE, AND M. TRAJTENBERG (2001): "The NBER Patent Citations Data File: Lessons, Insights and Methodological Tools," Working Paper 8498, NBER. [963]
- HSIEH, C.-T., AND P. J. KLENOW (2009): "Misallocation and Manufacturing TFP in China and India," *Quarterly Journal of Economics*, 124, 1403–1448. [948]
- JOVANOVIC, B., AND G. M. MACDONALD (1994): "The Life Cycle of a Competitive Industry," *Journal of Political Economy*, 102, 322–347. [947]
- KAPLOW, L. (2005): "The Value of a Statistical Life and the Coefficient of Relative Risk Aversion," *Journal of Risk and Uncertainty*, 31, 23–34. [968]
- KÖNIG, M. D., J. LORENZ, AND F. ZILIBOTTI (2012): "Innovation vs. Imitation and the Evolution of Productivity Distributions," Discussion Paper 8843, CEPR. [947]
- LAMOREAUX, N. R., AND K. L. SOKOLOFF (2003): "Intermediaries in the U.S. Market for Technologies, 1870–1920," in *Finance, Intermediaries, and Economic Development*, ed. by S. Engerman et al. Cambridge: Cambridge University Press, 209–246. [947]
- LUCAS, R. E., JR., AND B. MOLL (2014): "Knowledge Growth and the Allocation of Time," *Journal of Political Economy*, 122, 1–51. [947]
- PERLA, J., AND C. TONETTI (2014): "Equilibrium Imitation and Growth," *Journal of Political Economy*, 122, 52–76. [947]
- RESTUCCIA, D., AND R. ROGERSON (2008): "Policy Distortions and Aggregate Productivity With Heterogeneous Establishments," *Review of Economic Dynamics*, 11, 707–720. [948]
- ROMER, P. M. (1986): "Increasing Returns and Long-Run Growth," *Journal of Political Economy*, 94, 1002–1037. [947]
- SERRANO, C. J. (2010): "The Dynamics of the Transfer and Renewal of Patents," *RAND Journal of Economics*, 41, 686–708. [948]
- (2015): "Estimating the Gains From Trade in the Market for Patent Innovation: Evidence From the Transfer of Patents," Unpublished Manuscript, Universitat Pompeu Fabra. [948]
- ZUNIGA, M. P., AND D. GUELLEC (2009): "Who Licenses Out Patents and Why? Lessons From a Business Survey," Working Paper 2009/5, Directorate for Science, Technology and Industry, OECD. [979]

Dept. of Economics, University of Chicago, 1126 E. 59th Street, Chicago, IL 60637, U.S.A. and NBER; uakcigit@uchicago.edu,

Dept. of Economics, University of Toronto, 150 St. George St., Toronto, ON M5S 3G7, Canada; muratacelik@gmail.com,
and

*Dept. of Economics, University of Pennsylvania, Philadelphia, PA 19104,
U.S.A. and NBER; non-cotees@jeremygreenwood.net.*

Co-editor Liran Einav handled this manuscript.

Manuscript received December, 2013; final revision received September, 2015.